Technical Documentation

Crab Tracker

Noah Strong

Revision 1.0-WIP
May 22, 2018

Contents

1 Introduction

The Crab Tracker project was designed as a cost-effective means of remotely tracking crabs through acoustic signals. Small piezoelectric transmitters can be waterproofed and attached to crabs, and their intermittent signals can then be received by a set of four piezoelectric receivers configured in a square array.

Because this product was built with very few "off-the-shelf" components, it is somewhat complex and many of the finer details may be difficult for future collaborators to infer based on the existing documentation. This document aims to provide a technical overview of the project, including the rationale for some of the design choices, in hopes of giving the reader a deeper insight into the inner workings of the product.

2 Background and Overview

2.1 Motivation For The Project

While solutions already exist to aid in the tracking of aquatic animals, they are often prohibitively expensive without significant financial resources. The Crab Tracker product has been designed with cost and simplicity in mind, so a majority of the components, including both hardware and software, are custom-made. However, much of the product, especially the software, has been designed in a "modular" fashion, meaning that the various components are not tightly coupled. Simply put, future researchers should be able to exchange various components to better fit their needs with relative ease.

2.2 Description of Components

Our solution uses high-frequency audio broadcasts to transmit data about each crab to a central receiving station (likely mounted on a kayak or similar water vessel). At the receiving station we have four piezoelectric receivers, or "hydrophones," which are tuned to detect the signals that the crabs' transmitters broadcast. These hydrophones are arranged in a square shape with equidistant side lengths, and all are to be submerged an equal depth into the water.

As acoustic transmissions reach the hydrophone receivers, two distinct processes must be run. The first process is to detect incoming signals and label them with accurate timestamp values. Once this is done, software must run calculations using these timestamps to decode the unique identifiers of the broadcasting transmitter(s) and to determine the direction from which the sound(s) originated. Because the later task requires a mild amount of computation time, and because the former task requires low-latency monitoring of the input devices, these two tasks are to be performed on separate physical computers.

The first of the two devices, whose task is to timestamp incoming data, need only be a simple microcontroller. As of the current revision of this document, we have tested the system with an

Arduino Nano, a small, inexpensive device with a clock speed of 16MHz. Once broadcasts are detected and labeled with a timestamp, the device can send this data to the main computer that will run the calculations. We originally investigated using the popular Inter-Integrated Circuit (I^2C) protocol for this, but unfortunately I^2C can't be interrupted during a transfer, and therefore the device would be effectively frozen (and therefore unable to detect incoming data) any time a transfer was in progress. If any data were to come into the device during a transfer, its timestamp would be delayed and therefore inaccurate. Instead, we opted to use a similar protocol, Serial Peripheral Interface (SPI). The advantage of SPI is that data can be placed into a separate physical bus on the device when it is ready to be transmitted, and the other machine can fetch that data without affecting the detection and timestamp code.

The second device, which runs calculations and displays results to the user, is a slightly more powerful machine. We opted to use a Raspberry Pi Model 3B for this task, as these machines are inexpensive, reasonably powerful, well-supported, SPI-compatible, and easy to develop on. We also purchased a 7-inch touchscreen made specifically for the Pi as a means for the user to interact with the system in the field. The Pi is directly wired to the Arduino for SPI. See Figure 1 for an illustration of these components.
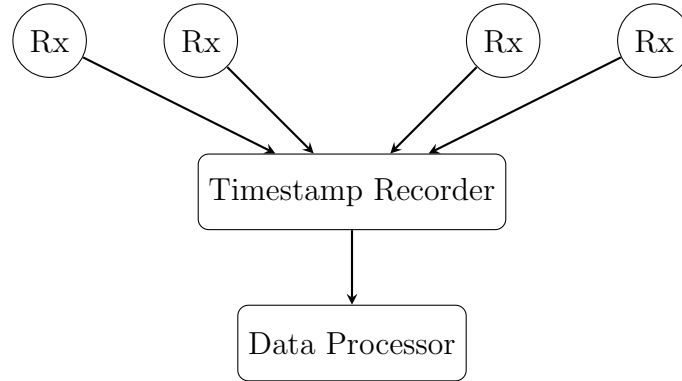

Figure 1: A simple diagram detailing the main components.

3   Electronics Hardware

While both the Arduino and the Raspberry Pi were off-the-shelf products, the transmitter and receiver hardware was all custom designed and fabricated for use in this project. Clearer specifications, including detailed diagrams, are detailed in other documents and therefore many details will not be repeated here.

3.1   Transmitters

The transmitters are small piezoelectric transducers that oscillate at an ultrasonic frequency. In our work, this frequency has been in the 40-60kHz range, depending on the resonant frequency of

the transducer itself. Though we have not established bounds on our range of potential frequencies, the actual frequency shouldn't matter so long as the receiving hardware is appropriately tuned to pick up this frequency (through filtering, hardware characteristics, and so on).

Each transmitter is attached to a small Printed Circuit Board (PCB) designed specifically for the project. The PCB can power the transducer by feeding it a current, effectively turning it off or on as needed. Each board will be programmed with a unique identifier (a numeric value). The ID determines the pattern of pulses that the transmitter should emit. For more details on this, see the iCRAB Transmission Protocol document.

3.2   Receivers

The receivers include small piezoelectric transducers that are able to detect vibrations in the same range as the transmitters. A preamplifier steps up the signal from the transducer before sending this signal along to a band-pass filter. The band-pass filter restricts the frequencies that the hardware "listens" to, meaning that a lot of irrelevant noise can be discarded. After being filtered, the signal is then amplified. A potentiometer on the board allows for variable amplification. Next, the signal is rectified (effectively producing the absolute value of the incoming sine wave). The low pass filter transforms the signal from a sinusoidal wave into a single pulse. Finally, the Schmitt Trigger acts as a variable threshold, preventing the rising or falling edge from being interpreted as a series of short rising and falling edges. See Figure 2 for details on this pipeline.
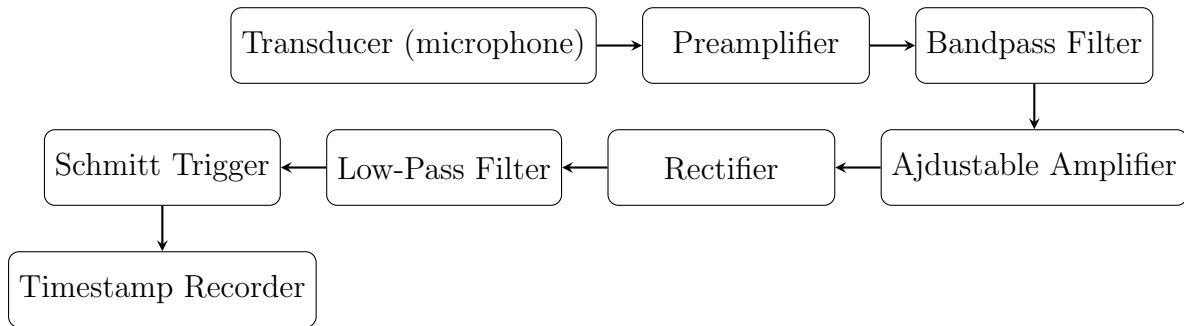
Figure 2: The receiver hardware pipeline.

The resulting digital signal is sent over a wire to the timestamp recorder detailed in section 4.1.

4   Computer Hardware

Signals detected through the receiving hardware described in Section 3 are delivered to more advanced microcontrollers and computers using common General Purpose Input Output (GPIO) pins. Two separate devices are used for the signal processing. The first, referred to in this document as the Timestamp Recorder (4.1), provides a low-latency means of detecting the time of an incoming signal. This device must be able to reliably detect when a signal arrives with

sub-millisecond precision. Once timestamped, the data can be transferred to the Data Processor (4.2). This device decodes the ID from each transmission and displays the crab's direction to the user.

4.1 Timestamp Recorder

The timestamp recorder reads a digital signal on a GPIO pin and records the "time" at which the signal changed from low to high or vice versa. In this case, the "time" need not have any meaningful relationship to "wall clock" time, so long as the timestamps have some meaning relative to each other. For example, the time associated with a state change could be measured in microseconds since the board reset. Since the software (described in Section 5.2) is primarily concerned with the time separation between events, the epoch is irrelevant. It is important, however, that the Timestamp Recorder and the Data Processor computers agree on the units used, because some of the later processing will be based on the speed of sound in water, measured in meters per second.

In addition to providing reliable timing measurements when input pin states change, the timestamp recorder must also be able to communicate the timestamped data to another device. Use of a non-blocking protocol such as SPI is necessary here so that the data transfers do not have a large negative impact on the device's ability to continuously detect changes.

4.2 Data Processor

The majority of the processing software exists on a separate device, referred to as the Data Processor. This machine reads timestamped data from the Timestamp Recorder (4.1) and stores it in its own local storage until it has enough data to run some calculations. This device is also the user's primary interface for interacting with the system in most use cases. Equipped with a touchscreen or other means of displaying data, the computer is able to display the unique identification numbers and relative locations of crabs as they are detected.

The software written for this project assumes that the Data Processor is a Linux-based machine with an SPI interface. Specifically, it has been tested on a Raspberry Pi 3B, which is a very suitable device given its low cost, built-in GPIO pins, and adequate processing power.

4.3 Wiring Diagrams

This section is currently blank. A diagram illustrating the connections between the receivers, Arduino, and Pi will be inserted here in future revisions of this document.

5 Software

Wherever possible, the software for this project was written in such a way that various elements may be altered or entirely replaced without requiring major changes to the rest of the codebase.

The hope is that it is flexible enough to be adapted for new scenarios and hardware setups in the future. That said, it has only been tested with the hardware detailed elsewhere in this document, so its portability and modularity is unproven.

In this section, we will examine the software that runs on the two devices described in Section 4.

5.1   Timestamp Recorder Software

The relevant source code can be found here: `src/arduino-src/arduino-src.ino`.

The Timestamp Recorder (4.1) software has two primary functions: (1) record timestamps of signal changes (i.e. when a pin goes high or low) and (2) transfer the timestamps and pin values to the Data Processor.

The `arduino-src.ino` file is heavily documented, so this document will step through the major sections without providing extensive detail about individual lines. Instead, we will examine in broad strokes the data formatting, data transfer, and input pin change detection.

5.1.1   Storage, Setup and Loop

Data is stored in a bounded buffer where each entry contains a timestamp value and the state of the pins at that timestamp (i.e. which pins were high and which were low). The timestamp is a 32-bit `unsigned long` and the pin values are all stored in a single `byte` where each bit corresponds to a single pin. The `pinvals` are shifted such that the lowest order bit corresponds to pin D3. A 1 in position 0 means that pin D3 is high; 0 means low. (Note that a few of the low-order pins on the board are reserved for special purposes, so the receivers are to be attached to pins D3 through D6.)

The `setup()` function sets the GPIO pins on Port D (pins 0-7) as input pins, meaning that another device (in this case, a receiver, as described in section 3) can drive the line high or low. The Timestamp Recorder can detect these changes when the pins are in input mode.

The `loop()` function runs indefinitely and on each iteration it can fill the SPI data register and/or detect and timestamp a change on one of the pins. We will examine both of these primary functions in more detail in the following two sections.

5.1.2   SPI Handling

The Arduino Nano has a built-in SPI bus that supports 8-bit data transfers. On a transfer, the Master node provides a clock pulse and on every pulse, one bit is shifted out of the data register (sent to the master) and one bit is shifted in (sent from the master). On the Arudino, the data register is referenced with the `SPDR` variable.

On each iteration of the loop, the Arduino checks to see if the "End of Transmission" flag is set, which indicates that a transfer has completed. If it is set, the board can fill the Data Register with

the next byte to be sent. The next time the master requests data, the newly-loaded byte will be shifted out to that device.

As mentioned, the built-in SPI hardware supports 8-bit data transfers. However, each segment of data that we wish to transfer contains a 32-bit timestamp and an 8-bit state, meaning that it takes 5 individual transfers to send a segment (or "block") of data. When the master node requests data, the slave sends the pin values on the first transfer, and then sends the timestamp in 8-bit chunks over the course of the next four transfers, starting with the lowest order byte and working up to the highest order byte. Note that this means both the master and slave nodes must be in sync with each other, as there is nothing to indicate what kind of data is sent in a single transfer. The two devices may get out of sync if one is interrupted between transfers, such as when the master's software is stopped unexpectedly. In this case, when the master's software restarts, it will expect to read pin values on the first transfer, even though the slave node may be prepared to send part of a timestamp.

To solve this issue, the SPI master node can send flags to the slave that will alter the slave's state or behavior. Certain values that the master can send will induce specific behaviors on the slave. For example, the `SPI_ECHO_REQUEST` flag, when sent, will cause the slave to send the `SPI_ECHO_RESPONSE` value on the next transfer. This is used when the master's program starts up in order to ensure that the two devices are correctly connected. The `SPI_RESET` flag will re-position the de-facto "read heads", meaning that it will send a new state starting with the pin values on the next transfer. This functionality is utilized to ensure that the two devices stay in sync when one of them restarts while the other runs.

5.1.3 Pin Change Detection and Timestamping

Pin change detection is comparatively straightforward. Each time the loop runs, the current state of the relevant pins (D3..D6) is compared with the state recorded from the previous iteration. If the two states differ, this indicates that at least one pin changed from low to high or vice versa. When this is the case, the device stores the new state in the global bounded buffer, along with the current time.

As of the current revision of the document, the timestamp is generated from the `TCNT0` clock, which increments every four milliseconds. This is an inadequate granularity for the purposes of this project, so future revisions of the code will instead read from a source that can generate ticks at a higher rate (for example, a clock cycle counter, or a timer set up to increment every few microseconds).

Note that no matter the units used, we should expect to store no more than 32 bits of information for any given timestamp. Therefore, $2^{32} \approx 4.3$ million is the highest value we can represent. Limited precision means that a source that increments frequently will overflow more rapidly than a slower source, which leads to an important trade-off. A 4-millisecond timer will overflow every $4 \times (4.3 \times 10^6)$ milliseconds, or roughly once every 4,772 hours. However, a timer that increments once per microsecond will overflow every 1.1 hours (or roughly 72 minutes). Currently, overflowing timer values are not handled by any part of the software; instead, any such data will automatically be thrown out. In other words, if a transmission is received while an overflow occurs, such that

some timestamps have extremely high values while others have extremely low (near-0) values, that transmission will be ignored. No attempt is made to account for the overflow, though to do so is well within the realm of possibility for future work.

5.2   Data Procesesor Software

The relevant source code can be found in multiple files in `src/pi-src/`.

5.2.1   Overview of Files

The primary functions of the various files are as follows:

- `main.cpp` – The main loop of the program. This file periodically polls the SPI slave device and then runs routines to process any data that is fetched.

- `data_collection.cpp` – When SPI data is read, the functions in this file parse and store the data in more meaningful structures. For example, when a rising and subsequent falling edge are detected on a given pin, a `ping` object is stored containing these timestamps. The data collection code also steps through the collected data to determine if enough pings have been collected for the direction algorithm to run.

- `direction.cpp` – Using an algorithm developed by one of the team's developers, Chloe Yugawa, this code determines the relative direction of a crab given the timestamps of its pings. That is, it performs a type of Time Difference On Arrival calculation based on ping times. It returns a direction relative to the front of the boat in polar coordinates in the form (r, θ, z).

- `spi.cpp` – This file facilitates and abstracts away the communication with the SPI slave device. Routines in this file can determine if the devices are properly connected and will also ensure that the master and slave devices are in sync with each other.

- `uid.cpp` – Each transmitter broadcasts in a specific pattern based on its unique identifier. The code in this file examines the duration of and delay between two pings in order to determine which transmitter created a given broadcast.

- `util.cpp` – Various helper methods are located here.

5.2.2   Processing and Structure

Each file contains extensive documentation, so the finer details about various functions and variable will not be duplicated in this document. Instead, we'll take a broad look at how these pieces fit together.

When data is read from the Timestamp Recorder, it is translated from raw pin states into something more meaningful, such as a "ping" (that is, a single broadcast as measured by the rising and

falling edges of a signal on a single pin). As pings are collected, the software looks to see if it has enough pings to run calculations. Specifically, it looks to see two pings on every hydrophone before it attempts to locate and report a crab. Recall from the requirements of the iCRAB Transmission Protocol that every transmission by a transmitter should be made up of two pings.

When enough pings have been amassed and a single crab's pings have been detected twice on each input line, the software can then run the direction algorithm with these values. If the direction algorithm is able to determine an approximate location for the signal*, the direction and associated ID can be reported to the user.

* Note that the direction algorithm will intentionally fail if it is given impossible data. That is, if the timestamps are too wildly separated temporally, it will assume that the measurements are not from the same transmission and will therefore cease to calculate a direction.

Glossary

GPIO General Purpose Input Output. 4–6

I²C Inter-Integrated Circuit. 3

iCRAB Transmission Protocol Our custom encoding protocol for transmitters. 4, 8

PCB Printed Circuit Board. 4

SPI Serial Peripheral Interface. 3, 5–8